

## Import the necessary libraries

```
In [1]: import yfinance as yf
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.stats.diagnostic import acorr_ljungbox
from sklearn.metrics import mean_squared_error, mean_absolute_error
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: # AAPL Stock data from 2011 to 2025
Stock_data = yf.download("AAPL", start="2015-01-01", end="2025-01-01")
```

YF.download() has changed argument auto\_adjust default to True

[\*\*\*\*\*100%\*\*\*\*\*] 1 of 1 completed

```
In [3]: # Display first 5 rows
Stock_data.head()
```

```
Out[3]:
```

	Price	Close	High	Low	Open	Volume
	Ticker	AAPL	AAPL	AAPL	AAPL	AAPL
	Date					
	2015-01-02	24.288586	24.757340	23.848711	24.746232	212818400
	2015-01-05	23.604334	24.137514	23.417722	24.057537	257142000
	2015-01-06	23.606550	23.866475	23.244431	23.668754	263188400
	2015-01-07	23.937571	24.037541	23.704304	23.815383	160423600
	2015-01-08	24.857302	24.915063	24.148616	24.266361	237458000

## Data Processessing

```
In [4]: Stock_close = Stock_data[['Close']].copy()

## the index changed to datetime
Stock_close.index = pd.to_datetime(Stock_close.index)

# Sort by date
Stock_close.sort_index(inplace=True)
```

```
In [5]: # Check for missing values
print("Missing values:", Stock_close.isna().sum())

# Drop or fill missing values
Stock_close.dropna(inplace=True)


# final structure
print(Stock_close.head())
print(Stock_close.info())
```

```
Missing values: Price  Ticker
Close  AAPL          0
dtype: int64
Price          Close
Ticker          AAPL
Date
2015-01-02  24.288586
2015-01-05  23.604334
2015-01-06  23.606550
2015-01-07  23.937571
2015-01-08  24.857302
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 2516 entries, 2015-01-02 to 2024-12-31
Data columns (total 1 columns):
#   Column          Non-Null Count  Dtype
---  -
0   (Close, AAPL)    2516 non-null   float64
dtypes: float64(1)
memory usage: 39.3 KB
None
```

# Exploratory Data Analysis(EDA)

## Line Plot of Closing Prices


```
In [6]: plt.figure(figsize=(12, 6))
plt.plot(Stock_close, label='AAPL Close Price')
plt.title('AAPL Stock Price Over Time')
plt.xlabel('Date')
plt.ylabel('Price ($)')
plt.legend()
plt.grid(True)
plt.show()
```



## Rolling Mean and Standard Deviation

```
In [7]: # 30-day rolling mean and std deviation
rolling_mean = Stock_close['Close'].rolling(window=30).mean()
rolling_std = Stock_close['Close'].rolling(window=30).std()

plt.figure(figsize=(12, 6))
plt.plot(Stock_close['Close'], label='Close Price', color='blue')
plt.plot(rolling_mean, label='30-Day Rolling Mean', color='red')
plt.plot(rolling_std, label='30-Day Rolling Std', color='green')
plt.title('AAPL Stock Price, Rolling Mean, and Rolling Std (30-Day)')
plt.xlabel('Date')
plt.ylabel('Price ($)')
plt.legend()
plt.grid(True)
plt.show()
```



### Augmented Dickey-Fuller (ADF) Test for Stationarity

```
In [8]: # Perform ADF test on the Close prices
result = adfuller(Stock_close['Close'])
print(f'ADF Statistic: {result[0]}')
print(f'p-value: {result[1]}')

# Interpretation:
if result[1] < 0.05:
    print("The series is likely stationary.")
else:
    print("The series is likely non-stationary.")
```


ADF Statistic: 0.7973433031053031  
p-value: 0.991593162707525  
The series is likely non-stationary.

### Stationarizing the Series

```
In [9]: # Differencing
# First difference (remove trend)
Stock_close['Differenced'] = Stock_close['Close'] - Stock_close['Close'].shift(1)

# Drop NaN created by differencing
Stock_close.dropna(inplace=True)

# Plot the differenced series
plt.figure(figsize=(12, 6))
plt.plot(Stock_close['Differenced'], label='Differenced Close Price')
plt.title('Differenced AAPL Stock Price')
plt.xlabel('Date')
plt.ylabel('Price Change ($)')
plt.legend()
plt.grid(True)
plt.show()
```



### ADF Test After Differencing

```
In [10]: # Perform ADF test again on the differenced series
result_diff = adfuller(Stock_close['Differenced'])
print(f'ADF Statistic (Differenced): {result_diff[0]}')
print(f'p-value (Differenced): {result_diff[1]}')

# Interpretation
if result_diff[1] < 0.05:
    print("The differenced series is stationary.")
else:
    print("The differenced series is still non-stationary.")
```

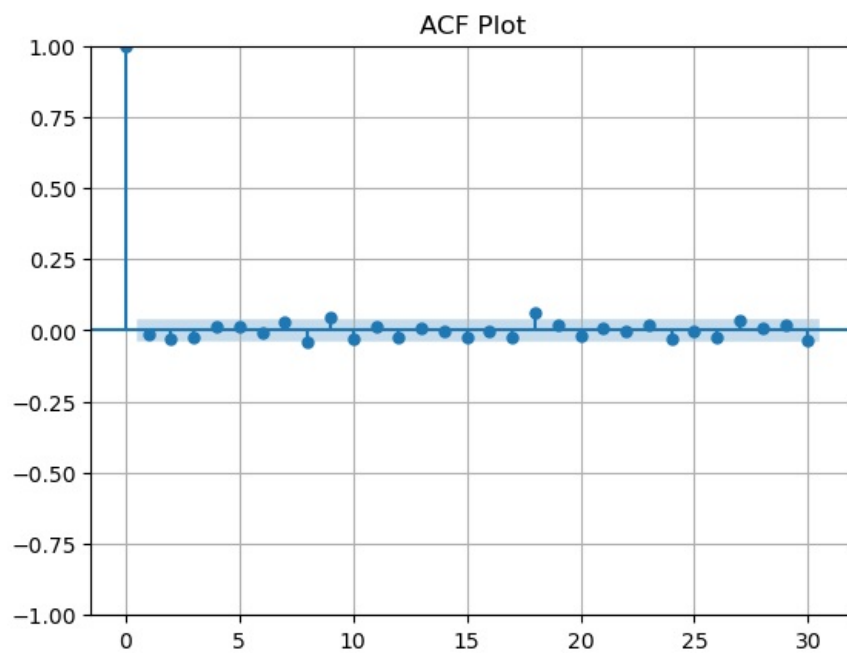
ADF Statistic (Differenced): -15.833236355123667  
p-value (Differenced): 9.886975652882048e-29  
The differenced series is stationary.

### Time Series Modelling (ARIMA)

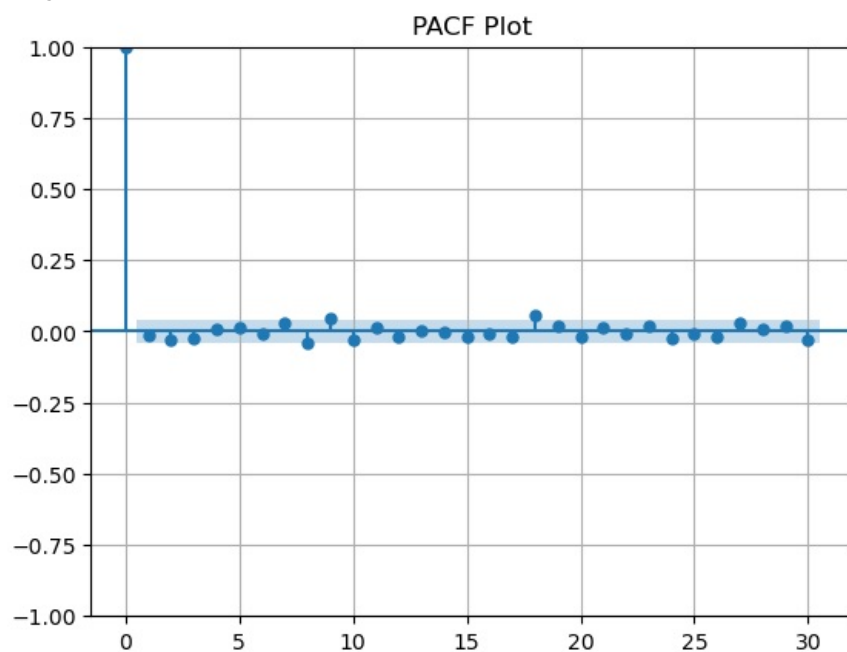
```
In [11]: plt.figure(figsize=(12, 5))
plot_acf(Stock_close['Differenced'], lags=30)
plt.title("ACF Plot")
plt.grid(True)
plt.show()

plt.figure(figsize=(12, 5))
plot_pacf(Stock_close['Differenced'], lags=30)
plt.title("PACF Plot")
plt.grid(True)
plt.show()
```

<Figure size 1200x500 with 0 Axes>



<Figure size 1200x500 with 0 Axes>



## Fit an ARIMA Model

```
In [13]: # Fit the model
model = ARIMA(Stock_close['Close'], order=(5, 1, 2))
model_fit = model.fit()

# Summary of the model
```

```
print(model_fit.summary())
```

#### SARIMAX Results

```
=====
Dep. Variable:          AAPL    No. Observations:          2515
Model:                ARIMA(5, 1, 2)    Log Likelihood          -5214.138
Date:                 Thu, 15 May 2025    AIC                  10444.277
Time:                 10:16:54    BIC                  10490.914
Sample:                0    HQIC                  10461.203
                        - 2515
Covariance Type:        opg
=====
              coef    std err          z      P>|z|      [0.025     0.975]
-----
ar.L1          0.9139     4.497      0.203     0.839     -7.901     9.728
ar.L2         -0.1956     3.272     -0.060     0.952     -6.609     6.217
ar.L3         -0.0031     0.077     -0.040     0.968     -0.154     0.148
ar.L4          0.0307     0.032     0.964     0.335     -0.032     0.093
ar.L5         -0.0040     0.147     -0.027     0.978     -0.291     0.283
ma.L1         -0.9258     4.498     -0.206     0.837     -9.741     7.889
ma.L2          0.1816     3.325     0.055     0.956     -6.336     6.699
sigma2         3.7069     0.054    69.265     0.000     3.602     3.812
=====
Ljung-Box (L1) (Q):                0.01    Jarque-Bera (JB):          3926.31
Prob(Q):                          0.91    Prob(JB):                0.00
Heteroskedasticity (H):            41.96    Skew:                    0.05
Prob(H) (two-sided):              0.00    Kurtosis:                9.12
=====
```

#### Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [15]: # Fit ARIMA(1,1,1) model
model = ARIMA(Stock_close['Close'], order=(1, 1, 1))
model_fit = model.fit()

# Summary
print(model_fit.summary())
```

#### SARIMAX Results


```
=====
Dep. Variable:          AAPL    No. Observations:          2515
Model:                ARIMA(1, 1, 1)    Log Likelihood          -5215.698
Date:                 Thu, 15 May 2025    AIC                  10437.395
Time:                 10:23:27    BIC                  10454.884
Sample:                0    HQIC                  10443.743
                        - 2515
Covariance Type:        opg
=====
              coef    std err          z      P>|z|      [0.025     0.975]
-----
ar.L1          0.5480     0.379      1.445     0.149     -0.195     1.292
ma.L1         -0.5678     0.375     -1.514     0.130     -1.303     0.167
sigma2         3.7115     0.053    70.204     0.000     3.608     3.815
=====
Ljung-Box (L1) (Q):                0.08    Jarque-Bera (JB):          3902.29
Prob(Q):                          0.77    Prob(JB):                0.00
Heteroskedasticity (H):            42.01    Skew:                    0.05
Prob(H) (two-sided):              0.00    Kurtosis:                9.10
=====
```

#### Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

#### Plot Fitted vs Actual

```
In [16]: # Plot predictions vs actual values
plt.figure(figsize=(12, 6))
plt.plot(Stock_close['Close'], label='Actual')
plt.plot(model_fit.fittedvalues, label='Fitted', color='red')
plt.title('ARIMA Fitted vs Actual AAPL Close Price')
plt.xlabel('Date')
plt.ylabel('Price ($)')
plt.legend()
plt.grid(True)
plt.show()
```




## Residual Diagnostic


```
In [17]: # Extract residuals
import seaborn as sns

residuals = model_fit.resid

# Plot residuals
plt.figure(figsize=(12, 4))
plt.plot(residuals)
plt.title('Residuals from ARIMA(1,1,1)')
plt.grid(True)
plt.show()

# Plot density of residuals
plt.figure(figsize=(12, 4))
sns.histplot(residuals, kde=True, bins=30)
plt.title('Distribution of Residuals')
plt.grid(True)
plt.show()
```





```
In [18]: # Test if residuals are white noise
ljung_box_result = acorr_ljungbox(residuals, lags=[10], return_df=True)
print(ljung_box_result)
```

	lb_stat	lb_pvalue
10	14.675135	0.144362


## Forecasting with ARIMA Model

```
In [19]: # Forecast 30 steps into the future
forecast_steps = 30
forecast_result = model_fit.get_forecast(steps=forecast_steps)

# Get prediction and confidence intervals
forecast = forecast_result.predicted_mean
conf_int = forecast_result.conf_int()

# Plot forecast with confidence intervals
plt.figure(figsize=(14, 6))
plt.plot(Stock_close.index, Stock_close['Close'], label='Historical', color='blue')
forecast_index = pd.date_range(start=Stock_close.index[-1], periods=forecast_steps + 1, freq='B')[1:]
plt.plot(forecast_index, forecast, label='Forecast', color='green')
plt.fill_between(forecast_index, conf_int.iloc[:, 0], conf_int.iloc[:, 1], color='lightgreen', alpha=0.4)
plt.title('30-Day Forecast of AAPL Stock Price')
plt.xlabel('Date')
plt.ylabel('Price')
plt.legend()
plt.grid(True)
plt.show()
```





## Model Evaluation: Train-Test Split & Forecast Accuracy

### Split the Data

```
In [20]: # Using the last 30 business days as test data
train_data = Stock_close['Close'][:-30]
test_data = Stock_close['Close'][-30:]
```

### Fit ARIMA on Training Data

```
In [21]: model = ARIMA(train_data, order=(1,1,1))
model_fit = model.fit()
```

### Forecast

```
In [22]: forecast_result = model_fit.forecast(steps=30)
```

### Evaluate Performance


```
In [23]: rmse = mean_squared_error(test_data, forecast_result, squared=False)
mae = mean_absolute_error(test_data, forecast_result)
mape = np.mean(np.abs((test_data - forecast_result) / test_data.replace(0, np.nan))) * 100

print(f'RMSE: {rmse:.2f}')
print(f'MAE: {mae:.2f}')
print(f'MAPE: {mape:.2f}%')
```

RMSE: 21.10  
MAE: 18.88  
MAPE: nan%

### Visualize Actual vs Forecast

```
In [24]: plt.figure(figsize=(14, 6))
plt.plot(test_data.index, test_data, label='Actual', color='blue')
plt.plot(test_data.index, forecast_result, label='Forecast', color='red')
plt.title('ARIMA Model: Actual vs Forecasted Prices')
plt.xlabel('Date')
plt.ylabel('Price')
plt.legend()
plt.grid(True)
plt.show()
```



In [ ]:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js